

From Aesop to AI-sop: A Guide to the Future of Storytelling with AI-Assisted Animated Audiobooks

By Polybios

July 2023

No More Dreaming, It's Time to Actualize Your Vision!



Yes, that's Socrates creating his animated audiobook!

Table of Contents

| | |
|---|-----------|
| Introduction | 3 |
| • Brief introduction of the project..... | 3 |
| • The problems that AI tools have solved for creators..... | 3 |
| • The purpose and goals of the project..... | 3 |
| • Opportunities and challenges of AI in the creative field..... | 4 |
| Concept Development and Prototyping..... | 5 |
| Overview of the main tools used..... | 6 |
| Writing (LLM's)..... | 6 |
| • Outline and themes..... | 7 |
| • Editing your text..... | 10 |
| • Other Prompt Examples and Creative Strategies..... | 15 |
| Visual Art..... | 17 |
| • Intro to Midjourney..... | 18 |
| • The “Zeroing in” method..... | 18 |
| • Maintaining Consistency..... | 18 |
| • Eliminating “Hallucinations”..... | 20 |
| Audio Generators..... | 22 |
| • Elevenlabs..... | 22 |
| • PlayHT..... | 22 |
| Music and Sound Effects..... | 23 |
| • Finding and downloading..... | 23 |
| • Copyright issues..... | 23 |
| • Slight modifications..... | 23 |
| Video Editing..... | 23 |
| Conclusion..... | 24 |

Introduction:

This project was a two-week long experiment of amalgamating various AI technologies (LLM's, art and audio generators) with my passion for Greek history and philosophy. It represents a potential prototype of a new media format, which could be a promising bridge between the dwindling superficial sensationalism of Hollywood, and books that many find challenging to concentrate on and commit to in our fast-paced, image-saturated world. AI technologies genuinely hold the power to enhance our creative potential, freeing us to create without the confines of entrenched principles or the financial pressures that high-budget films frequently encounter. This concise guide seeks to offer a sweeping understanding and kindle creativity, guiding you on your journey.

In the pre-AI era, I aspired to craft something akin to this—narration paired with art to depict scenes. Upon my calculations, the cost was prohibitive, which often deterred me. The expense for quality voice narration of 10k words would minimally be \$500 and could skyrocket to \$5k, especially considering hiring six voice actors for my project's dialogues. Commissioning an artist to create 50 illustrations would cost a minimum of \$1k, (I used over 300 in my project). Hiring someone to refine the text, polish the writing, and provide critical feedback would run another \$500. In essence, I would be spending a minimum of \$2k, not to mention the time and energy required to source these professionals, communicate my vision, and negotiate terms. Furthermore, I'd be bound to accept whatever they produced.

The intention of this project, as outlined in the epilogue, was to tackle a multitude of present-day issues by drawing parallels with events that transpired during Plato's and Aristotle's time. However (and hopefully) not appropriating them in the past but merely rediscovering them. Additionally, the project aimed to portray the spirit of Socrates in a more captivating manner, with the hope of inspiring individuals to explore Platonic philosophy in a less daunting, more accessible way. Ultimately, my goal was to streamline dialogues and illuminate the following themes (if you have seen the video with the epilogue then skip the following three paragraphs):

Plato's famous Allegory of the Cave, which symbolizes the pursuit of truth, is often used today to dissect modern issues and forms the basis of the movie "The Matrix". Further explored are the risks associated with declining educational standards, as well as the challenges of democracy, including the rise of demagogues who in turn are influenced by foreign powers. Alongside this, it examines the danger of powerful institutions dictating morality. Lastly, it raises the issue of philosophical relativism of truth, which appears to underpin many disquieting social patterns of today.

In the Sophists, we observe an unhealthy emphasis on argumentation, even at the expense of truth, a trait prevalent in today's society and educational system, as evident in its fascination with debate clubs. Instead of rigidly defending our positions, we could nurture 'dialectic clubs,' aimed at approaching the truth and finding common ground.

Plato's teachings are often misinterpreted due to cultural appropriation and erroneous translations by western academia. For instance, the concept of "agape" or teacher-student love in ancient Greece, is incorrectly interpreted as pederasty. Yet, in reality, the Greeks recognized multiple forms of love. This specific one reflects the mutual

respect, reciprocal learning and the necessary closeness between teacher and student to foster an environment conducive to dialectic methodology. Of course, such a notion clashes with our modern assembly-line education system, designed in the 18th century to mass-produce obedient workers, where students are passive learners who are not encouraged to think critically or creatively. An education system that is progressively undermined by ideological agendas, and overshadowed by the might of the internet and artificial intelligence.

Opportunities and challenges of AI in the creative field:

AI provides a springboard for human creativity. It offers tremendous opportunities in the creative sector by automating repetitive tasks. It allows for personalization of content, enhancing user engagement and experience. AI serves not just as a tool, but also as a creative partner, contributing novel ideas and perspectives, as well as offering constructive criticism when requested. Moreover, it democratizes the creative landscape, making advanced tools more accessible to individuals without formal training or background. Finally, AI empowers us to achieve our fullest creative potential by overcoming barriers to entry, assisting us in articulating our ideas in the most effective manner, and enabling us to express them artistically in ways many of us couldn't before. In other words, it enables us to become the best version of ourselves.

One of the challenges of using AI tools for this specific type of storytelling is ensuring consistency throughout your content. For example, making sure that the images follow a similar theme, artistic style, lighting, background and level of detail. Another challenge is grasping the true capacities of these tools and determining the best prompting strategies so you can get the most out of them. I created this journal partly to share the techniques I employed to address such challenges.

On a broader level, as more and more creatives start using AI to augment their work (or even if they don't use it at all), they may face accusations that their creations are not genuinely theirs. This can lead to talented individuals being undervalued due to heightened competition and questions surrounding the authenticity of their work. The Vice President of Microsoft stated that AI will write your first draft. However, from my experience, if you want to create something truly unique and creative, driven by your passion, AI can't quite compete yet. AI is a tool that enhances you. People don't want mediocrity; they want greatness, and thus far, greatness can only come from human creators.

The final challenge I foresee is the dependence on technology. Over-reliance on AI could potentially result in the loss of certain human skills and the under-appreciation of human creativity. For instance, in writing, we still need to understand the fundamentals of good writing to draft the first version and have the judgement to refine what the AI has produced. We also need to know how to demand better versions, as you'll see later on, I often asked AI to provide 10 different versions of a sentence. From there, I'd choose the top three and integrate them.

One last note, if you may, is that when creating art, it's essential to proceed with respect and caution and not use it as a tool for political appropriation, a trend increasingly prevalent today.

This is something I've strived to avoid. In "The Republic," Plato warns of the dangers of poets and art. He offers the analogy of a painter who creates a portrait of a man. Bystanders admire the painting for its realism, some even claiming it to be more lifelike than the man himself. However, when a shoemaker passes by, he notices that the painting has an inaccuracy in the depiction of the man's shoes. In our image-saturated society, where we stay indoors and reduce our interactions with the world, consuming vast amounts of content, there's a risk of art instilling a skewed perception of reality. Over years and decades, this false sense of reality could leave us disillusioned, disappointed and lead us astray.

Concept Development:

They say innovation and creativity often occurs when diverse technologies and concepts are **combined** in fresh, effective, and captivating ways. The real potential of these AI technologies is when we begin to weave them together in innovative and creative ways, whether for art, software, business, or social initiatives. However, there's no replacement for inspiration, passion, persistence and the formation of teams and communities to bring these plans to fruition.

Perceptions of originality and authenticity can often mislead us. It's more beneficial to find a topic that sparks your passion—preferably one you're already well-versed in—and then accumulate insights from a range of sources to remix and reconstruct in unique (and simple) ways. Culture is always dead and reborn. It becomes obsolete when it loses its utility and is resurrected when it regains relevance. Countless pieces of content lie dormant in antiquated libraries, untouched and unread. You have the potential to breathe new life into these forgotten works, repackaging them in an appealing manner suitable for our digital, image-driven age. Your creation doesn't need to be limited to an animated audiobook; it could manifest as an online course, a training video, a component of a larger project, or something else.

The generation of initial ideas, the refinement of your concept, and the ultimate vision of your project largely rest in your hands. However, some guidance on these areas will be covered in the writing phase of this guide

Prototyping:

Prototyping is a crucial aspect of any creative endeavor. It's beneficial to start by creating a simple prototype of your project. For instance, focus on the initial stages of the project or a small segment, disregarding the finer details of the entire project. This will give you a sense of the overall direction and help identify any limitations you might face. For my own project, I produced a rough cut of the first scene (3 minutes), using a basic script and the free version of the art generator and a basic free audio subscription. I assembled these elements together, despite their initial rawness, to understand the overall process. This gave me valuable insights and the confidence to progress. (As a matter of fact, the whole video I made was a prototype of a much larger book or series of books about Plato's life. In retrospect it was rather hastily made and as I say in my opening sentence it was an experiment of AI tools, and an indulgence of creative inclination.)

If you're planning a 15-minute video, consider creating a rough draft/version of the first minute. Don't invest all your initial time into writing or audio production trying to make it perfect,

because the creative process often involves iterating back and forth between different stages and seeing in practice if something looks good. This concept, known as a Minimum Viable Product (MVP) in the field of innovation, is expected to be imperfect at the beginning.

It's also worth noting that you shouldn't feel obligated to read this entire document at the outset. If you read everything, you might feel overwhelmed. Start small: write a few words in a document, next, utilize GPT by providing it with the context and requesting it to generate ideas, solicit criticism, and even assist in conducting research on your topic. Do a quick polish, grab a free version of the audio and art software, and play around with them. Throw them all into the video editor and create your first prototype. The journey of creation is iterative, and taking manageable steps will help you stay focused and motivated. Seeing the first prototype can be both satisfying and energizing! It's a moment that sparks a multitude of ideas for improvement and refinement, truly kickstarting the creative process!

Overview of the main tools used:

For a video editor I used Shotcut as its free, easy to use and YouTube is filled with "how to videos" for it. [GPT-4](#) was used as the main LLM. While GPT-4 is more creative, I sometimes used Bard, whenever I felt stuck with GPT or when I wanted to double check or get a second opinion on something, but most importantly I used Bard for finding these AI tools, (as it has better access to the internet). The LLM's are the foundation of everything. They help you refine/polish your writing, brainstorm ideas, strategize, do research, fact checking, find the tools you need, and you create prompts for those other tools. For the visual art I used [Midjourney](#), as it seems to be the best one so far. For the narration audio I used [ElevenLabs](#), and for the conversational audios I used [PlayHT](#).

Writing (with LLMs):

As someone with dyslexia, which I would define as your brain working faster than you can produce words, I've often felt restrained in expressing myself. It's as if there's a funnel restricting the flow of my ideas. But now, with AI tools like ChatGPT, I no longer worry about phrasing or even typos. I write as fast as I can, then paste my text into ChatGPT, providing it with context, tone, and the scope of what I'm writing about and use my most common prompt: "**rephrase this**". Then for sentences that really matter to me or that the ones the AI didn't phrase as I wanted, I ask GPT to express that sentence in ten different versions, and then I select or merge the versions I like best.

Furthermore, an AI language model like GPT will only be as good as the input and instructions you provide. The more questions you ask from it, the more perspectives you request, the more precise, demanding and multidimensional your instructions are, the more valuable it becomes. However, quantity of your prompts is just as important as their quality as you'll see later on.

It's important to note that when using AI for idea generation or text creation, only about 5-10% of the output might be truly valuable, however some responses might simply spark further inspiration. GPT is a tool to enhance your writing, not to do the work for you. Most of the time, the results won't align perfectly with what you envisioned, but that's all part of the process: sifting through, creating variety, and selecting the best.

Enough theory for now; let's dive straight into practical examples of how I prompted the AI. My intention is to stir your creative thinking and inspire you to devise your own innovative prompts. To keep this concise, I won't display the AI's responses unless it is used to illustrate something. The responses of GPT are left in its signature gray color, the prompts I used are indented, or highlighted in green.

Outline and Themes:

By no means is this a writing guide on how to write well. I'm not even good at that. I'm just giving you ways to use the AI. (However, you can have GPT give you teaching lesson and tell you all about it and how your favorite writers write!)

When preparing the outline, ensure you provide as much information and context as possible, presenting your ideas in a raw manner. Engage with the AI, asking it to help you develop the outline collaboratively. Before requesting the AI to create the outline based on your information, encourage it to ask you questions about the story and context, as LLMs do not possess implied knowledge. This interaction will aid the AI in better understanding the narrative and allow it to offer ideas for consideration. During the process, express what aspects you like and dislike.

Similarly, when exploring themes, involve the AI by requesting it to provide more than 20 themes to choose from. This collaborative approach will lead to a more enriched and comprehensive outline and themes for your writing project.

(Note that the prompts you will see for creative writing are from the second chapter of my video, the historical fiction. The first chapter is all history, with the only exception being a minor exaggeration concerning the Marduk priests. As for the Delphic priests there are multiple sources talking of ties with the Persians, including the famed historian Herodotus)

Here is an initial prompt about the themes:

"The story will explore the effects of sophists on society. Demagogues will be shown manipulating the crowds and making them unstable. The people will be shown not thinking for themselves and being manipulated. The Delphic oracle will also be shown manipulating society and stirring up civil war. Persian agents will be shown influencing Athenian politics. Plato will struggle with his ideas, but he will also struggle with how to act when his city is decaying. Platonic dialogues will be condensed into 5-7-minute versions. The value of friendship will be shown, as Plato develops strong friendships. Geopolitical and strategic aspects, as well as secret intelligence games, will be shown. Plato's growth and maturation will be shown, as well as his internal and external struggles. The importance of education will be shown.

In summary, the themes will be education, philosophy, tyranny, the failures of democracy, foreign influence, theocratic manipulation of society, friendship, service to the state, and so on. These themes will be developed gradually and may not all be covered in the same chapters.

Do you have any comments, questions, or suggestions? try to think like a great novelist and encapsulate the works of Steven Pressfield, Fyodor Dostoevsky, Paulo Coelho, and Dan Brown (The Da Vinci Code)."

Then I asked it to give me some suggestions.

After that I gave it this prompt for creating an outline of the first chapter of the historical fiction (second chapter of the video):

"The first chapter of the novel will introduce the protagonist, Plato, a young student of philosophy. The Peloponnesian War has been waging for 26 years, and Athens is not faring well. Plato witnesses the decay and corruption in Athenian society, and he begins to question the value of sophists and demagogues. He meets Socrates, and the two of them begin to build a relationship. Plato recognizes Socrates's value, and he begins to grow as a philosopher.

At some point in the chapter, Plato will attend the Pythian Games and win the wrestling championship. A priest from the temple will approach him and offer him a position of power. The priest is a bad guy, and he is secretly influencing Athenian politics. Plato will have to decide whether or not to accept the offer, and he will also have to decide whether or not he wants to become a politician. The chapter will end with Plato still unsure of what he wants to do with his life. He has a naïve belief in himself to become a politician and help Athens, but he also believes that he is not ready. He will have to decide what is more important to him: his own personal beliefs or the good of Athens.

The chapter could open with a scene of Plato wrestling in the gymnasium. This would introduce him as a physical as well as intellectual being. The chapter could also include a scene of Plato witnessing a philosophical discussion in the gymnasium. This would show his interest in philosophy and his desire to learn more. The chapter could explore the relationship between Plato and Socrates. This would show how Plato's views on philosophy and politics are shaped by Socrates. The chapter could also include a scene of Plato meeting the priest from the temple. This would introduce the priest as a villain and set up the conflict between Plato and the priest."

(Note that I have lightly edited and spell-checked all the prompts I have provided here. I write quickly, which often leads to spelling errors and gibberish, so I wanted to make sure the prompts are clear and easy to understand)

After receiving the initial rough outline of the first chapter of the historical fiction part from GPT, I proceeded to provide a prompt that addressed that previous outline. I asked GPT to ask questions and present alternatives to collaboratively build a better first chapter. This approach is crucial to the creative process.

Chat GPT: "sure! It's always good to explore different options. Let's consider a few alternatives to various parts of the current outline, and I'll also ask some questions to guide your thought process:

1. **Starting Point:** While the current plan is to start with the description of Athens during the war, an alternative could be to start with Plato's wrestling match. This could provide a more immediate and gripping introduction. *Question: Which do you feel would more effectively draw the reader in?*
 2. **Introduction of Plato:** We've planned for Plato's introduction in the gymnasium. Alternatively, we could introduce him in an educational setting, emphasizing his intellectual prowess right from the start. *Question: Which context do you feel would better serve the character and the story?*
 3. **The Meno Dialogue:** In the current outline, Meno is introduced as a Persian agent who is subtly influencing Athenian politics. An alternative approach could be to show him as a seemingly benign figure initially and gradually reveal his true affiliations. *Question: Would a more gradual reveal add to the suspense and intrigue of the story?*
- ... 4,5"

Then I let it know what I liked and asked it to incorporate those in a new outline.

After that, I asked it again if it has any other scenes to suggest. This is how you take advantage of the AI, it is inexhaustible, the more you prompt it for ideas the more likely it is that you will end up with a good one!

Certainly, here are a few additional scenes we could consider:...

And then I continued with another prompt:

"I want you to make Xenophon to be Plato's best buddy, and also put Alcibiades in the story, showing him as someone who's really keen to get ahead. I also want to include a part where Plato heads to a fancy brothel one night, maybe with his friends, maybe by himself. There, he runs into a young, quick-witted, and gorgeous hetaira, and they have a little chat. When he's leaving her place, he crosses paths with Alcibiades for a moment and they exchange a few words. Alcibiades comes off as flashy but also pretty smart, and Plato feels a bit intimidated around him."

Next prompt:

"Are there any general things that we could improve for the first chapter? what would a great writer like Stephen Pressfield and Dan Brown suggest we do?"

ChatGPT

Stephen Pressfield and Dan Brown are both fantastic authors, albeit with different styles. However, they do share some fundamental principles of good writing that could be applied to your chapter:

1. **Clarify the narrative voice:** Both Pressfield and Brown are known for their strong narrative voices. Establish a consistent point of view and narrative tone.
 2. **Balance between exposition and action:** Both authors excel at balancing descriptive passages with action to maintain reader interest. Consider if there are moments in the chapter that could be enlivened with more specific, concrete details or moments of action or suspense.
- ... 3. 4. 5. 6. 7.

(Follow up) ok how do these translate into our chapter 1?

ChatGPT

Absolutely, here's how you can apply these principles to the first chapter of your book:

1. **Clarify the narrative voice:** The voice seems factual and somewhat academic in the current text. Consider whether a more engaging, narrative voice could make the content more accessible. You could achieve this by occasionally speaking directly to the reader, asking rhetorical questions, or employing a more conversational tone.
... 2, 3, 4, 5, 6, 7.

After that I did a series of prompts trying to instruct it to consider all the principles we discussed and write me the first scene. That was not effective at all.

GPT can struggle with processing and implementing abstract or broad instructions effectively (such as: write me this scene) , particularly when tasked with making nuanced changes while preserving the original structure of a text. Additionally, AI often generates shorter responses than requested and can occasionally veer away from the provided context. To harness its power, it's advisable to break down your instructions into smaller, more specific tasks, and feed the AI relevant context as needed. It requires careful human guidance and fine-tuning to maximize its potential.

After a series of disappointing prompts and after me becoming increasingly precise in my prompts, I found myself gaining a not very acceptable output. Regardless, I ended up extracting some valuable segments, merging them into the original text. In my estimation, this process improved the quality of my writing by around 20%, while also achieving a comparable condensation. **Brevity is crucial in writing**, so when asking GPT to refine your text, also ask it for a 25% reduction, and noting that key ideas should be maintained (most of the times it will do it by itself). This might require multiple tries. Now, let's move onto editing your text.

Editing your text:

Prompt: "Can we employ the seven principles (that great writers employ) we've discussed to enhance the opening scene of the book's second part? This scene needs to be enthralling and serve as an effective introduction to where the dialogues begin. Please provide suggestions for improvement and potential additions. Additionally, pose questions to me about my preferences for editing and expanding this section."

Note: it's always beneficial to prompt the AI to ask you questions about the section before it starts making edits or suggestions. This approach not only helps the AI understand the context better, but also provides you with a variety of innovative ideas to choose from.

ChatGPT

"First, let's focus on the introductory principle:

Opening the scene with action: The scene currently opens with a dialogue, it may benefit from a more active introduction. You could consider opening with Plato engaging in some form of physical activity or philosophical contemplation.

Now let's discuss the other principles:

Creating dynamic dialogue: The dialogue in this scene is functional and conveys necessary information about the characters and their relationships. To add more depth, it could be helpful to insert some subtext or indirect communication. This may make the dialogue more engaging and realistic.

.... (more topics followed; they are not included so as not to make this document too big)

Can you provide more information on the following points:

1. What kind of person is Plato at this point in his life? What is he passionate about, and what are his internal conflicts?
2. How does Plato's relationship with his mother and uncle impact his decisions and beliefs?
3. Could you provide some details about the setting? Where is this conversation taking place?
4. Are there any specific philosophical or political concepts that you'd like to introduce or explore in this scene?
5. How would you like to build suspense or tension in this scene?
6. What major events or decisions are looming over Plato at this moment?

Answering these questions will help us further revise and enhance the scene.

In its response right above, GPT suggested the concept of incorporating a philosophical contemplation as we introduce Plato standing by a window. I found this idea intriguing and subsequently requested it to present ten different approaches to realize this concept. This method can be analogous to panning for gold. Initially, you sift through a lot of material (context/problem or your first draft), asking the AI to generate a variety of ideas and pose related questions. This initial sifting (step A) is like shaking the pan - you're looking for the promising nuggets (ideas) among the gravel. Once you spot potential 'nuggets', you carefully examine them (request the AI to generate ten versions of the chosen idea(s)), similar to inspecting each gold piece for its worth. Finally, from these versions, you select the ones that shine the brightest (ideas you like the most) and then request the AI to convert the ideas into actual text,

Prompt: "Generate ten distinct concepts for philosophical contemplations that we could use to initiate our opening scene. Keep in mind Plato's philosophies, but also remember that he is just 20 years old at this point in the story and is still maturing in his thought process. Maybe just present some of the initial contemplations and questions that might have sparked his famous ideas."

ChatGPT:

1. Suspense or tension in this scene?
2. What major events or decisions are looming over Plato at this moment?
3. Certainly, here are some philosophical contemplations that a young Plato might engage with:
4. Nature of the Physical World: As Plato is observing the sunrise over Athens, he might be contemplating on the inherent beauty of nature and how it contrasts with the conflict and turmoil of the human world. This could hint at his later theory of Forms and Ideas, where the world of Ideas represents perfection, in contrast to the physical world.

5. Justice and Equality: While observing the Agora, he might reflect on the inherent inequality and injustice in Athenian society. He could question why some people have more privileges and others suffer, which would later form the basis of his thoughts on social justice and equality in his ideal Republic.
6. The Role of Education: ...
7. Beauty and Aesthetics: ...
8. Individual vs Society: ...
9. Virtue and Morality: ...
10. Truth and Knowledge: ...
11. Politics and Leadership: ...
12. The Soul and Immortality: ...

I selected three of these ideas and requested multiple renditions of each from the AI. Afterwards, I integrated these into a cohesive paragraph, blending them with my own thoughts and ideas. I even asked it to make the amalgamation.

Here is another example of the 'panning for gold' methodology:

1st prompt: "I want to end the dialogue with a joke. I'm aware that Socrates' wife, Xanthippe, was a vibrant and lively character, often portrayed as arguing with Socrates. Give me 10 topics of a what that joke could be about."

GPT: Sure, here are 10 topics around which a joke involving Socrates and his wife, Xanthippe, could be crafted: 1,2,3 ...10

2nd prompt: "Okay I like idea number 2, give me 15 ideas on what that joke could be about."

GPT: Absolutely, here are 15 ideas for a joke about Xanthippe's temper and how Socrates deals with it:

3rd prompt: "I like number 3, 5 and 6. give me 20 variations of those jokes"

GPT: Sure thing! Here are 15 variations of jokes 3, 5, and 6:

4th prompt: "I liked number 4 but it was a little bit cheesy and too long, can you make it shorter? Give me 20 variations of it."

As you can see this is a much better process than asking it just write you a joke. Instead, we gave it a direction ("make him say a joke about his wife"), we asked for volume to have better chances of stumbling upon something good, and then when we saw something promising we gave it further directions ("it's too cheesy, make it shorter") asking it for further variations on that specific one, to end up getting its best possible version.

Let's now see some other type of prompts:

Prompt: "Let's proceed to the second scene, applying the seven principles we've previously outlined. Please provide suggestions on improvements, potential additions, and pose questions. Here is the current draft: ("text was inserted here")"

ChatGPT

GPT gave me 7 questions and 5 suggestions. I actually copied and pasted those questions to a word document. And next to each one I added a brief answer and then re pasted the original text with my answers (highlighted as green) attached back to GPT. Here is how it looked:

1. How would you characterize Plato and Xenophon's relationship? What are the dynamics between the two characters? They share a strong bond of friendship, built upon mutual respect, and recognize each other's merits.
2. Are there any particular aspects of their characters you'd like to showcase in this scene? Show that Plato is a very good wrestler but also humble. For Xenophon I don't know. Can you make something up based on his historical character?
3. Could you provide some context for their discussion about "the enemy within"? I'm not sure about that, maybe we can remove it. I just wanted to show that they are cognizant of the decline of the city's morals.
4. Do you want to incorporate any philosophical ideas or debates into their conversation? Maybe, but don't focus on philosophical ideas, focus introducing the character and relationship of the two.
5. How does the gymnasium setting influence their interactions and the overall mood of the scene? It doesn't.
6. Can you tell me more about the Sophists who are arriving? What is their role in this scene and how do Plato and Xenophon perceive them? They have no role in the following scene, it is just meant to be transition to the next one.
7. Is this the first time Plato is encountering Socrates, or is he already a familiar figure to them? This is the first time.

In terms of improvements, I would suggest:

1. Describing the physical setting of the gymnasium and the atmosphere more explicitly to immerse the reader in the scene. Yes do that
2. Highlighting the contrast between the physical wrestling match and the philosophical discussion to underscore the tension between the physical and metaphysical worlds. Not it's too cliché.
3. Adding more variation and nuance to the characters' voices and reactions. For instance, Plato could express a mix of curiosity and skepticism towards the Sophists, while Xenophon could show more straightforward enthusiasm. Yes I guess it could make it more realistic. But don't make him look stupid.
4. Developing the idea of the "enemy within" further. This could refer to individual moral failings, societal decay, or the dangers of ignorance and complacency. It would be interesting to have Plato and Xenophon debate this concept more deeply. Yes but don't go too deep on this, they are just in their twenties and have just finished working out.
5. Building up to the arrival of the Sophists and Socrates in a way that heightens anticipation and reflects their significance in Plato's life. For instance, the friends could notice a change in the gymnasium's atmosphere as these figures approach, and their reactions could reveal their different attitudes towards them. Yes I like that.

I didn't like it's response, so I prompted it again:

"That was awful. Add lines, don't just alter my own. Be creative. You are also being very hyperbolic. Also, don't give a narrative explanation for each line of conversation. Just reduce the narration and write more dialogue. Be great!"

I didn't like it, so I prompted it once again, however instead of asking it to add text I asked for a critique of what we have so far, so it can give a direction:

"Be critical of our scene so far. It seems quite shallow, not showing the depth of our characters. What would you suggest we do?"

Once again, the outcome was unsatisfactory, so I decided to retain the original text. However, I discovered a new type of prompt that proved to be highly beneficial in many instances. I began asking GPT to provide a list of ideas, allowing me to selectively choose from the options presented. This approach has been instrumental in improving the quality and variety of responses I receive:

"Let's jump back to scene 2. Don't rewrite the scene, just give me some jokes that show they are good friends. And maybe alter the part about "the enemy within". It seems too cliché. Give me 10 joke ideas and 10 alternatives for the talk about "the enemy within". And then, if you have any other good ideas on how to make this a better scene (I pasted the whole scene here so it had it as a reference).

Note: The more general your instructions are, such as "rewrite the scene based on these instructions," the poorer the outcome you'll likely receive. Instead, strive for specificity. Ask it to generate 10 different versions of a certain sentence, solicit it for joke ideas, or request specific change suggestions. Since it can only write up to 1000 words (usually it will just generate 400 words regardless of what you tell it), it's advisable to request for the additional parts instead of a complete rewrite. This method encourages it to concentrate on your exact requirements:

Prompt: "let's head to the next scene. It's not a big deal. We'll put Plato in Delphi, showing he won a wrestling match. This can help us show a bit of ancient Greek sports traditions too. Have you got any ideas to make this better? (I inserted the draft)"

(I had gone over this scene a few times already but it's always good to re visit it again and try to improve it)

GPT: In terms of improvements, I would suggest:

The scene you've outlined is quite engaging and colorful, painting a vivid picture of the ancient Pythian Games. However, there are a few areas where further enhancements could be made to deepen the narrative and immerse the reader further into the setting and action.

Here are a few suggestions:

1. Detail the Atmosphere: Emphasize the hustle and bustle of the Games, the anticipation in the air, the sounds, the smells, and the color.
2. Delve into Plato's State of Mind: ...
3. Expand on the Match: ...
4. The Aftermath: ...
5. The Cultural Significance: ...

I followed up with this prompt:

"Okay, forget (number) 1. Give me one or two sentences for 2. Maybe a couple of sentences for 3. Yes, talk about 5, and also give 4 a try. Just give me the additional sentences, I will add them to the original if I deem fit. Don't rewrite the other parts. "

Note: the last clarification of just giving the additional sentences and not rewriting the whole thing was crucial.

For the next scene I gave it my draft and prompted it this:

"Alright, let's delve into some possible additions. Don't rewrite the entire scene, but provide me with the supplemental pieces so I can incorporate them if they fit well. Could there be a mention of Socrates and how Plato has become his disciple? Provide me with 15 scenarios depicting what this conversation could entail. Bear in mind, Socrates will ultimately defy the oracle's plans. However, historically, the oracle deemed him the wisest Greek, which complicates the narrative. How should we navigate and reconcile this contradiction?"

As I've mentioned before when I want to phrase a sentence the best possible way, I will usually provide that single sentence to the AI and ask it to rewrite it in 10 different ways, with or without specific guidance. Then, I choose the top four versions/sentences that appeal to me, and I merge the most compelling aspects of each one. I recommend doing this by copying and pasting those four sentences to a word document, then highlighting the words or phrases you find most effective in each and then merging them. I find it's generally best to do this integration yourself, especially if you're meticulous, but you can experiment with asking GPT to perform this integration. My guiding principle is to write the initial draft and make the definitive decisions on the final detailed version. In this way, you have control over the beginning and the end of the process.

Other sample prompts:

"Say this sentence but in a funnier way, give 10 different versions". "Say this using an analogy", "Be more brief/ poetic" etc. "Make her speak more eloquent". "give me some ideas, what would a person like that say?"

Well, why am I even trying to write this by myself? Let's end the writing section with some other strategies that GPT gave me that I completely agree with (I know that's cringe) :

1. **Get Various Takes:** For critical sentences or paragraphs, ask GPT to write it in multiple ways to explore different phrasing and tone possibilities.
2. **Character Dialogue:** If you're stuck with dialogue, ask GPT to generate multiple versions of the conversation based on character personality and story context.
3. **Character Development:** Ask GPT to generate background details, quirks, or traits for your characters to make them more three-dimensional.
4. **Exploring Genres:** Give GPT the same plot but ask it to write the scene in different genre styles (mystery, romance, thriller, etc.).
5. **Show, Don't Tell:** Ask GPT to show emotions and actions rather than telling them outright. For example, ask it to describe a character's sadness without using the word "sad."
6. **Scene Description:** Struggling with setting descriptions? Ask GPT to describe a scene based on the specifics you provide.
7. **Perspective Shifts:** Ask GPT to write a scene from the perspective of different characters to help you explore different viewpoints.
8. **Conflict Creation:** Ask GPT to suggest conflicts or plot twists that could enhance your story.
9. **Resolve Plot Holes:** Describe a plot hole or narrative inconsistency to GPT and ask for potential solutions.
10. **Opening Lines:** Ask GPT for a range of opening lines to kickstart your story or new chapters.
11. **Closing Lines:** Similar to openings, GPT can provide a range of powerful closing lines for chapters or the novel itself.
12. **Metaphors and Similes:** Ask GPT to generate metaphors or similes to vividly describe characters, emotions, or situations.
13. **Symbolism and Foreshadowing:** Ask GPT to suggest symbolic elements or foreshadowing techniques that can be embedded into your narrative.
14. **Title Suggestions:** Share a brief synopsis of your story with GPT and ask it to generate a list of potential titles.
15. **Scene Expansion:** If a scene feels too short or lacking in detail, ask GPT to help expand it.
16. **Scene Reduction:** Conversely, if a scene is too verbose or slow, ask GPT for help in condensing it.

17. Flashbacks or Flashforwards: Ask GPT to write a flashback or flashforward scene based on your current narrative.
18. Mood Setting: Ask GPT to generate descriptions that set a particular mood for a scene.
19. Critique and Revision: After writing a chapter or scene, ask GPT for a critique or suggestions for improvement.
20. Remember, while GPT can provide a myriad of ideas, ultimately, your creative intuition should guide the storytelling process.

Visual Art (Midjourney):

I made the artwork using Midjourney, a popular AI service. It requires a Discord account, which is simple to set up if you don't have one already. Midjourney offers a free version where your images are mixed with the images of others, but for a comprehensive project, I'd recommend the monthly-billed standard plan. This allows your work to be neatly organized and easily accessible.

Midjourney operates by inputting `"/imagine:"`, followed by a description of the image you want. the secret to obtaining a satisfying result is a mix of quantity and diversity—the more prompts you provide, and the more varied these descriptions are, the more likely it is you get what you were looking for. After you enter a prompt, Midjourney generates an image containing four variations. If you input "a cat in a hat," for instance, you might see the cat wearing different kinds of hats and in various poses.

Below is an example when I prompted the Midjourney bot to `"/imagine: Socrates in his ancient clothes in front of computer screens video editing. blue background"`:



As you can see, it returns four versions of the image. If I want to enlarge and maybe eventually use any of these four, I click U1 for the first image top left and U4 if I want the fourth image on the bottom right. If I want to see a variation of the first one, I click V1. I can keep going through the rabbit hole of seeing variations of variations indefinitely until I get something I like.

The process of finding/creating the right image is about bombarding the AI with a variety of prompts. You'll see many images, but they might not be exactly what you want.

In terms of prompts, mix up your approach. Use a combination of 5-10 different prompts of your own and 5-10 prompts that you asked GPT to create and then paste on Midjourney. Pay attention to 'focus words'—specific keywords that can guide the AI towards your desired output. These can be different in each situation. For instance, when I was trying to create an image of a navy of many Greek ships, I noticed that using grand words like "massive" or "impressive" lead the AI to create larger-than-life/non-realistic ships. If that's not what you want, adjust your 'focus words' accordingly. In the case of a ship, for instance, you might instruct the AI to imagine a "medium-sized ship". Once you've fine-tuned your approach and explored various images, you'll notice that certain prompts yield visuals closer to your preference. Proceed by taking these promising prompts, making slight adjustments, and resubmitting them as new prompts.

To recap this refinement process, you have two main strategies. The first approach involves identifying images you appreciate and instructing the AI to generate further variations of them (with the V1, V2.. method showed above), and then request additional variations of these initial variations. The second strategy revolves around pinpointing images you favor, understanding the prompts that birthed them, and subtly modifying these prompts. After that, sift through all the images (20-50), select your top 2-5 favorites, save them, open them in full screen, and then determine the one that truly captures your vision.

You can think of this as an artillery barrage; quantity can be important. You're essentially sweeping across a wide target area, trying to pinpoint the perfect impact point. By launching around 10-20 different prompts for the same scene, you are more likely to end up hitting near something you wanted. Once you notice that you are close to it, you can adjust your barrage's trajectory to precisely strike the intended target. Sometimes, the results from your prompts will align closer to the target, while other times, the prompts of GPT will score the hit.

To ensure a relative consistency between your scenes, it might be beneficial to produce not only 20-50 images for the scene you're working on, but also generate the same number of images for the following scene. This way, you can compare and select images that exhibit similar characteristics or color schemes, thus fostering a harmonious flow between scenes. Additionally, think about instructing the AI to depict characters showing their back. This strategy prevents viewers from seeing facial details that might be difficult to replicate in future scenes.

Note: you should not right-click on an image and save it. You first want to open it in another tab which has better quality, and then save it. Furthermore, in retrospect, I regret not adjusting the dimensions of all my images to fit the ideal YouTube video format, which is 1280x720 pixels (a 16:9 aspect ratio). You can achieve this by using an online image editor like Canva (<https://www.canva.com/>), where you can create a free account (there might be easier methods). By doing so, your images will fully occupy the screen, preventing the display of a square image with a black background on the sides...

Generally speaking, after many prompts for a specific scene, you will realize the limits of the AI. For instance, when you feed it various prompts featuring “a rich home interior, in classical Athens, a sitting area, and two people conversing”, the AI typically generates around five core themes and then variations of those. However, it's important to remember that the true constraint often lies not with the AI, but with your imagination and articulation of ideas. There's always potential to broaden the scope beyond the standard range, although generally, you will hit that boundary eventually.

Reaching these limits however, is actually beneficial—it's like pinpointing the core of your vision. It allows you to hone in further and make minute alterations to existing prompts. Such as adding to your prompts something like “from the two people conversing, one is a middle-aged man with a red clothing and the other is a young man in his twenties and his back is facing use”. And thus, your prompts become increasingly precise. Finding these limits is akin to cornering a wild animal – it's now much easier to tame it.

You can always learn more about prompting AI art generators by watching tutorials on YouTube, but I personally avoided this. My concern was that it would lead me down a rabbit hole of videos, flooding my mind with unnecessary information. My approach was to plunge into the project head-on, learning from experience and searching for specific details as needed.

Generate prompts from you narration (original text):

If you are unsure on how to prompt midjourney on a certain scene, you can copy and paste the text form you novel that involves that scene on GPT and ask it to produce 10 different prompts for this scene. You will usually have to specify what you want (as the AI does not know the context or other things that are implied), so if your sentence includes “Plato”, make sure you say: “Plato in his 20's”. You can also copy and paste the whole paragraph of the narration to Midjourney. Or paste parts of the paragraph and gradually add more in each new prompt to create a variety, as the example bellow:

This is the full narration from of the original text:

“The Pythian Games hosted by the temple of Delphi are in full swing. The grandeur of Delphi is awe-inspiring – an oasis were trees and marble blend seamlessly. the ornate temples, the tall statues, the imposing amphitheater. At the heart of it all lies the Temple of Apollo, a marvel of architecture and the spiritual nucleus of the Games. The city, mystical and sacred, seems to pulsate with the energy of thousands of spectators and athletes who have gathered to honor the sun god Apollo.”

1st prompt (to Midjourney): “/imagine: the Pythian Games hosted by the temple of Delphi are in full swing”

2nd prompt (add one more sentence): “/imagine: the Pythian Games hosted by the temple of Delphi are in full swing. The grandeur of Delphi is awe-inspiring – an oasis were trees and marble blend seamlessly.”

3rd prompt (add one more sentence): “/imagine: the Pythian Games hosted by the temple of Delphi are in full swing. The grandeur of Delphi is awe-inspiring – an oasis where trees and marble blend seamlessly. the ornate temples, the tall statues, the imposing amphitheater.”

And so on. Or any variation of the sentences you like.

Eliminating hallucinations:

AI art generators can sometimes generate images that are out of the ordinary, or even “hallucinations” as they are called. These might be something like a skyscraper in a landscape view of ancient Athens, or people having three hands or what I like to call “ghosts” (weird, out of context human looking creatures). However, you might really like that image and want to keep it, so in this case you can manually erase that hallucination. If the hallucination is at the side of the picture, you can simply crop it but that is not always ideal. For erasing hallucinations, I used Microsoft Paint, its free and already installed in your PC (I’m sure Mac has something similar).


To subtly erase elements in an image without leaving noticeable traces, select the “Color Picker” tool, zoom in, click on the color in the vicinity of the object you wish to erase, then switch to the “Brush” tool, set the brush to a size that fits your need, and carefully paint over the object with the selected color, effectively making it blend with the surrounding area and creating the illusion that it has been seamlessly removed from the image. Consider selecting colors from both sides of the object—using the “Color Picker” tool, pick one color from the area on the object’s right and another from its left, then use these two shades to paint over the object, creating a more natural transition that helps the erased portion blend even better with its surroundings.

Behold, my mildly successful before-and-after artistic ventures:

Socrates’s Daemon?



Voldemort lingering on the right side of the screen:



On the left, a flying candle distracting Socrates from his deep thoughts:



A mountain-sized mausoleum that if existed would be the 8th wonder of the world:

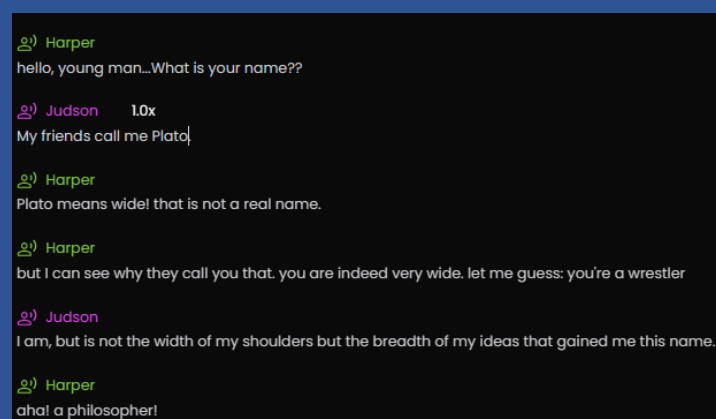


Audio Generators:

I used Elevenlabs (<https://elevenlabs.io/>) for the narration voice. This site offers an extensive selection of realistic voices to choose from. My recommendation when using Elevenlabs is to generate small portions of your text, around 1-2k characters at a time. This way, if you're not satisfied with a specific section (such as its tone), you can easily re-generate it without consuming a significant amount of your credits.

When working with dialogues, I found PlayHT (<https://play.ht/>) to be the most user-friendly platform. To begin, I recommend sampling all of their available voices. You can use filters to select specific characteristics, such as a 'middle-aged male with a British accent'. Keep a record of the voices you prefer, and then proceed with the dialogue creation process.

Creating dialogues on PlayHT is a straightforward process. Simply create a new line, select your preferred voice, and then add your text.



When it comes to audio generators (both Elevenlabs and PlayHt), there aren't many guidelines to follow. However, it's essential to remember that each audio creation from a specific text may slightly differ from one generation to another. Similar to all AI tools we've examined, the key is to experiment with multiple iterations for optimal results.

The use of punctuation, exclamation points, and other symbols will also affect the variation in the audio. These symbols can be interpreted by the AI as instructions to change the tone or emphasis of the audio. For example, an exclamation point might be interpreted as an instruction to make the audio sound more excited. When creating dialogues using PlayHT, you may find it necessary to run certain lines several times, (with and without making alterations to the text) to achieve the desired result. Here's a helpful tip: Instead of erasing a certain line every time you want to try a new one, consider creating multiple lines, one after the other. After generating the audio for all of them, listen to each version and then delete the ones you don't like. This way, you can easily compare different variations and choose the best-sounding dialogue lines for your project. Like this:

Music and Sound Effects:

I sourced my music and sound effects directly from YouTube. To convert these into MP3 format, I utilized a 'YouTube to MP3' site, where all I had to do was paste the video's URL to initiate the download process. (However, be mindful that such sites can sometimes be slightly dubious, peppered with numerous 'download now' buttons which are often misleading and are merely links to advertisements. Therefore, identifying the legitimate download button might take some tries). You can then drag and drop the tune into your video editor's audio track. It's essential to credit the artists you used in your YouTube description, failing to do so may result in your video being removed. You should look for music tracks that are copyright-free; they typically indicate this in the description or title with phrases like 'copyright-free' or 'no copyright.' Each artist has different crediting requirements, often outlined in the descriptions of their videos. One artist even requested an email from those who used his tracks, which I complied with.

When it comes to finding the right song, I had already known most of them, but you can also search on YouTube using the term 'copyright-free' along with your desired type of music, e.g., 'inspiring cinematic background music copyright-free' or 'ambient background music copyright-free' or 'melody piano music copyright-free' etc. The same applies for sound effects, e.g., 'rooster crowing sound effect,' 'cricket at night sound effect,' 'people laughing sound effect.' Using sound effects like a rooster crowing or crickets chirping can subtly signal to your audience a transition to another scene.

Video Editing:

As previously stated, I used Shotcut (<https://shotcut.org/>) video editor because it's free, easy to use, and common, allowing you to easily find YouTube videos for any questions you might have. For instance, you could search for "how to lower the volume on a track," "how to create transitions between videos," or "how to add text on Shotcut". Surprisingly, the editing process was more time-consuming and labor-intensive than I anticipated. However, this is a crucial stage where all elements come together and make the final product. Harmonizing the audio, music, and imagery at the right moment is essential for creating an appealing result. The volume of the music should be adjusted so it's neither too loud nor too soft, especially if it's playing in the background while dialogues or narration are taking place. I suggest using varying volume levels and testing both with and without headphones, selecting a level that is a balance between the two.

In terms of the songs, it was often necessary to edit and select specific parts (e.g., some songs may have a softer, melodic part that transitions into a louder segment). To maintain the more intense part but with a lower volume, you can select the track, position the selection pin (which indicates where your timeline is), and press Ctrl + S. This will split your track precisely where you want it (this method also applies to videos), allowing you to adjust the volume of the chosen segment of the audio.

If you wish to omit the louder, more intense melody and extend the softer melody, you can again cut it, then copy and paste to create a duplicate of the needed part. Try to find a point where the end of the first section is similar to the beginning of the next one (or edit them accordingly) and create a transition between these parts."

Conclusion:

These methodologies and tools merely serve as a launching pad; there is an array of additional prompting techniques and resources (such as text-to-video), that you can explore. I encourage you to delve into these possibilities and create an animated audio that surpasses mine! The key takeaways for me were, firstly, the significance of meticulous attention to detail, and secondly, the realization that iterative improvements can vastly enhance the final product beyond your initial expectations!

With the newfound understanding of the power of creative prompting, you're on the path to better results. Being a brilliant writer isn't a prerequisite; what matters is your ability to recognize that a scene may be lacking something, like humor or depth. Even if you can't pinpoint what's missing, you can engage GPT to identify the gap or request an imagined critique from a renowned artist. The process is fluid: shifting from a focused perspective to a higher-level view, zooming in and out on the issue, and then zooming back in. As you continuously expand and contract your viewpoint, creating a variety of options to choose from, keep revisiting your work with a critical eye and iterating on your creation, you're setting yourself up for success. But remember not to overwhelm yourself - it all starts with crafting your prototype!

My hope is that we can use these emerging technologies to good effect, to help us realize our potential and lead meaningful lives by actualizing our creativity.

May you find eudaimonia in your creative pursuits!

